



Achieving Faster Boot Time With Linux

Christopher Hallinan
Technical Marketing Engineer

mentor.com/embedded



©2013 Mentor Graphics Corp. Company Confidential

Android is a trademark of Google Inc. Use of this trademark is subject to Google Permissions.

Linux is the registered trademark of Linus Torvalds in the U.S. and other countries.

Agenda

- System Definitions
- Hardware Considerations
- Typical Embedded Linux System
 - Boot Sequence
- Measurement Methods
- Optimization Techniques
- Summary/Q&A

Define what "Fast Boot" Means

- Fast Boot is not a single technology or "product"
- Many techniques are architecture or platform dependent
- Product/application defines "fast boot"
- Your system requirements determine limits
 - CAN Bus in 50 ms
 - Rear view camera video in 2 seconds?
 - Partial HMI in 3 seconds?
 - Full multimedia plus networking in 4 seconds?
 - Do you need Secure Boot?

What takes so much time?

- Power/Clock Stabilization
 - usually negligible but should be considered
- Low Level CPU Initialization - ~ 100 ms
 - Bootloader (often multi-stage, ie secure boot)
- Loading images (kernel, u-boot, rootfs, dtb)
 - Usually these images live on NOR or NAND Flash
 - Even a small reduced kernel can be 1-2 MB (compressed)
 - Often stored compressed (kernel)
- Subsystem (Driver) initialization
- Mounting a root file system
- Userland – System Utilities and Applications

Typical Embedded Linux System

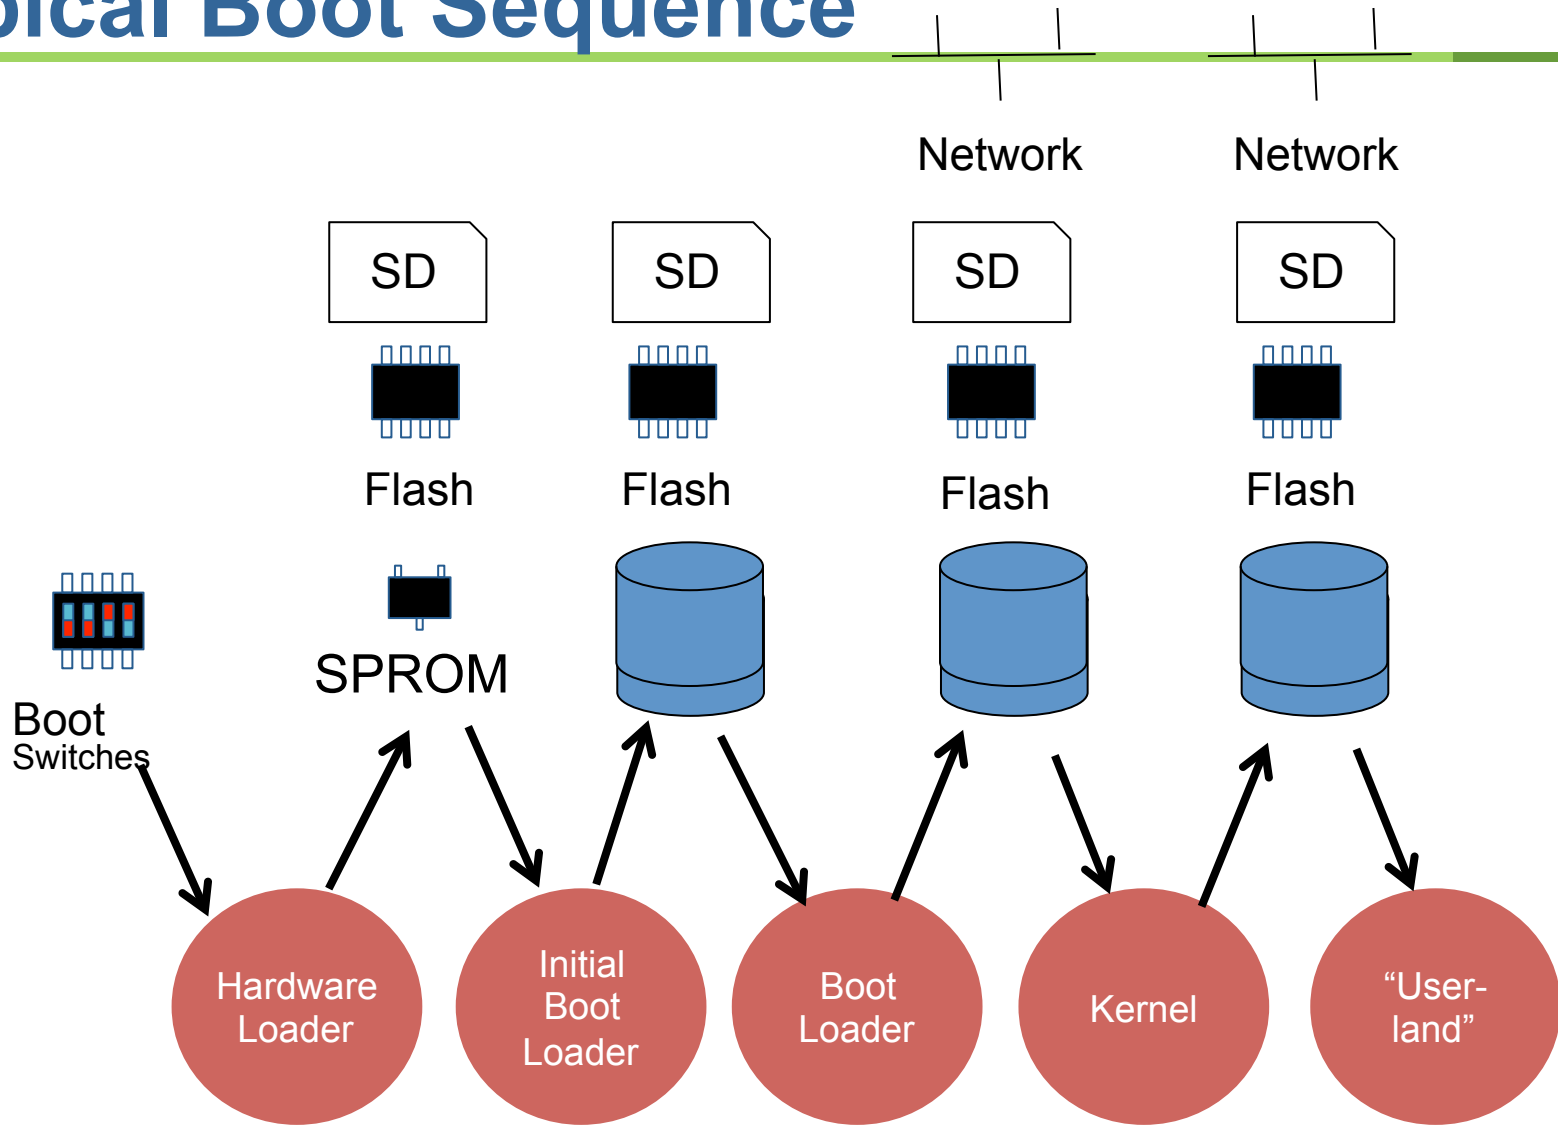
- Freescale i.MX6 SabreLite:
 - Quad-Core ARM® Cortex A9 1 GHz
 - Yocto core-image-sato (~90 MB rootfs)
 - Booting from Class 10 micro-SD
- Stopwatch analysis:

From Power On to:	Time (seconds)
Kernel FB logo	5
Userland psplash	13
Full Mobile Desktop	23

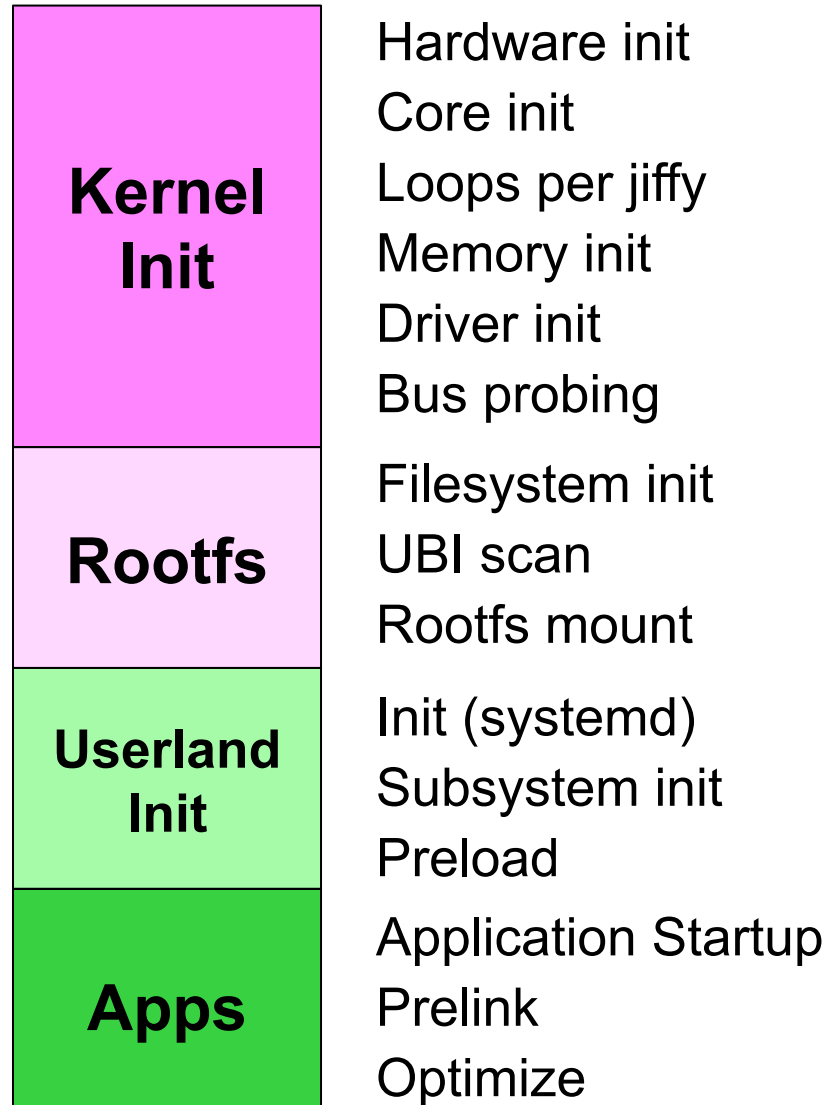
Hardware Considerations

- Hardware architecture makes a difference
 - Goes beyond just clock speeds, etc.
- Power and Clock stabilization should be very fast
- Design choices should support fast boot requirements
- Examples:
 - Loading u-boot and kernel from SPI NOR takes substantially longer than from parallel NOR or NAND
 - NAND flash or SD/MMC requires early software overhead but may be faster overall
 - Is your bootloader enabling caches early?

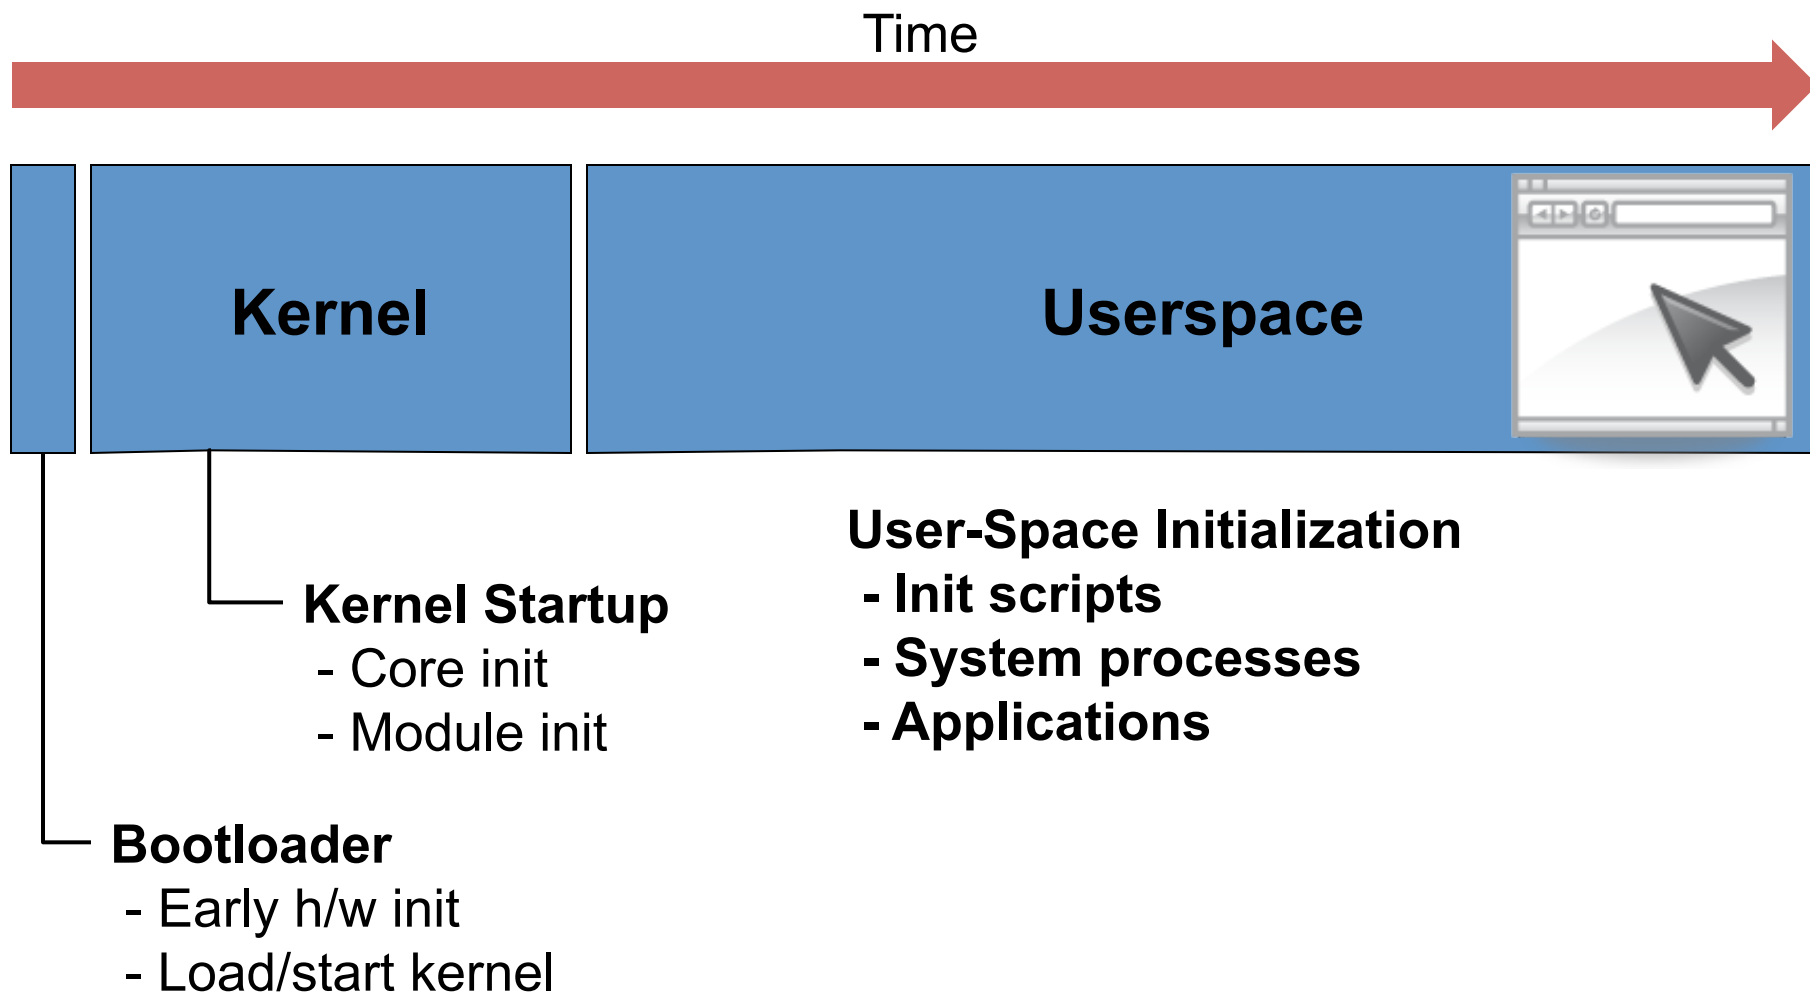
Typical Boot Sequence



Typical Linux Boot



The Bootup Phases (Relative Phase Lengths)



Splash Screens – Make it look like it's booted!

- Indicates system is active, but still booting
- A splash screen can take place:
 - In the bootloader
 - Will get splash up sooner, but...
 - Usually require some porting to bootloader
 - May delay early kernel functionality
 - In the kernel
 - After initialization of the framebuffer driver
 - Early user-space init (psplash)
 - Before system apps initialized



Profiling and Measurement Tools

Boot Time Measurement Methods

- Techniques depend on context:
- Bootloader profiling often requires custom tools
- Several tools are helpful for Linux kernel profiling
 - Some are very easy to use (CONFIG_PRINTK_TIME, etc)
- Userland tools
 - Many tools to choose from
 - Some easy to use, some require investment in learning
- Some portions require custom techniques

Profiling U-Boot

- No clever "out-of-the-box" tools
- Whatever you do here will be custom
- First order: something similar to `CONFIG_PRINTK_TIME`
 - Enable timestamp values on each line of console output
 - May require architecture-specific timer
 - Can also use "grabserial"
- Second order:
 - Custom time checks around suspect areas
 - Hardware tracing on supported processors

Some Popular Measurement Tools

- CONFIG_PRINTK_TIME
 - ftrace
 - Bootchart (userspace)
 - SystemTap
 - LTTng
- initcall_debug
 - oprofile (both)
 - perf
 - strace (userspace)
 - uptime
 - ...

Timing Printk

- A simple method to put a timestamp on every printk
- Useful to identify lengthy init operations
- Activation (use one of the following):
 - Compile kernel with: `CONFIG_PRINTK_TIMES=y` (in Kernel Hacking)
 - Use “`printk.time`” on kernel command line
 - Both of these methods can be used during kernel boot
 - Or, dynamically in a run-time system (as root):
 - `echo "Y" >/sys/module/printk/parameters/time`
 - Obviously only useful after kernel is booted

Using ftrace for profiling


- Requires some "investment" (learning curve)
- Analyze and debug latency and performance issues inside the kernel
- Profile execution time of functions, events, more
- A “framework” of several assorted trace and event utilities
 - Lots of filters and options for fine tuning your measurement
 - Over 800 events in 26 categories (syscalls, sched, irq, module, etc)
- Controlled via `/sys/kernel/debug/tracing`
- Can be enabled on kernel command line to facilitate early boot profiling. Ex:
 - `ftrace=function_graph tracing_thresh=5000`
- See `.../Documentation/trace/ftrace.txt` for more information

Bootchart

- Bootchart

- Very useful for correlating CPU utilization with initialization process
- Helps identify opportunities for parallel init
- Limited to userland initialization profiling
- Very easy to use: `init=<path-to-systemd-bootchart>`
- Also works with systemd

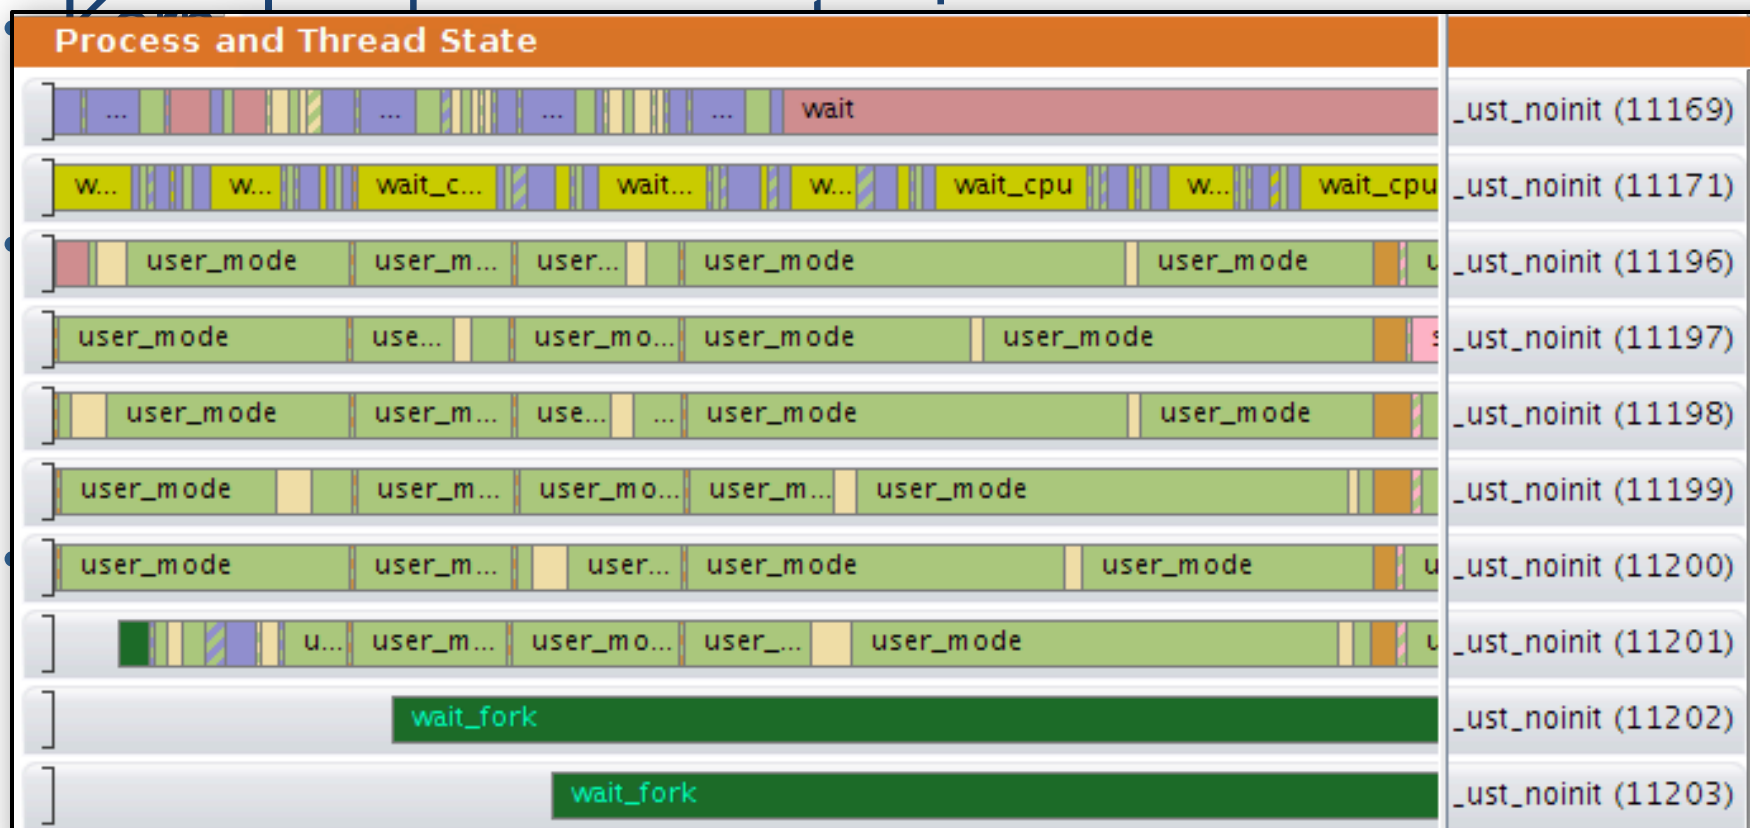
Typical Bootchart Output – i.mx6 startup



SystemTap

- Powerful instrumentation framework for tracing, profiling and evaluating kernel behavior
- Command line interface and scripting language
- Ideal for complex tasks that require live analysis without having to recompile the kernel
- Primarily designed for users with in-depth kernel knowledge and experience
- Can provide insight that no other tool can
- Significant learning curve
- Plenty of documentation on-line
- May not be available on every platform

LTTng (Linux Trace Toolkit next gen)

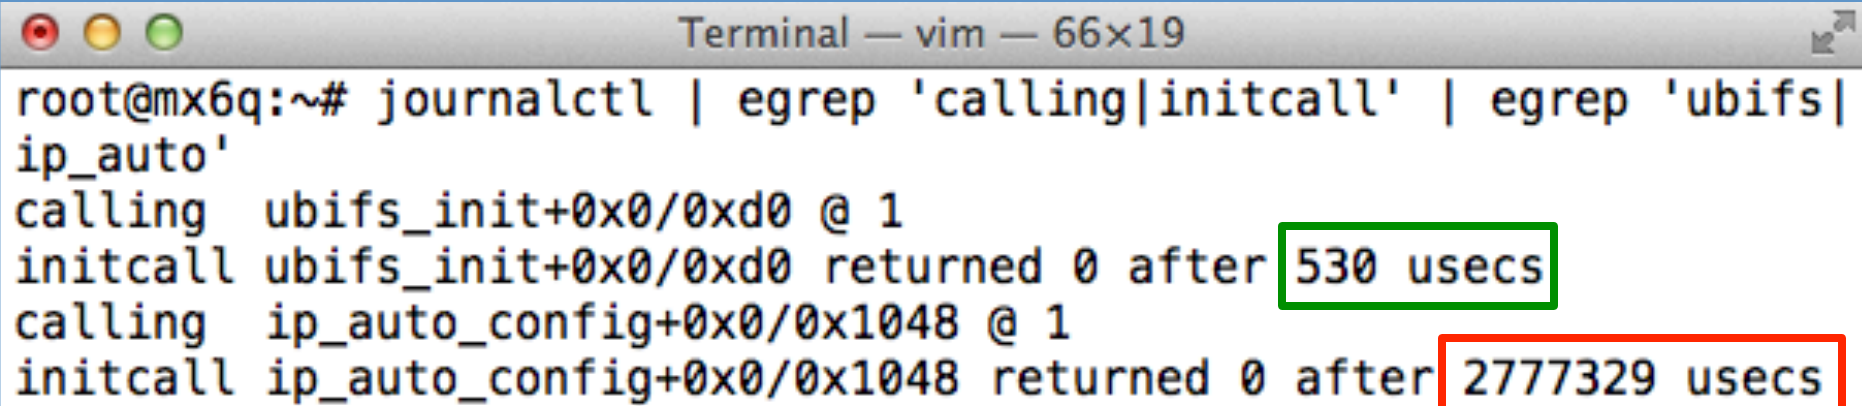


- For most modern kernels, LTTng is easy to integrate into your kernel
 - As always, YMMV depending on version and architecture

Using initcall_debug

- Driver initialization calls (initcalls) spend considerable time on kernel bootup
- A kernel flag enables initcall information during startup
 - On the command line, add “initcall_debug=1”
- Notes:
 - Increase the printk log buffer size in kernel config:
 - LOG_BUF_SHIFT=18 (256KB)
 - Enable CONFIG_KALLSYMS (to see function names)
- After booting info is found in bootlog (dmesg)
- May identify:
 - Opportunities for parallelism
 - Opportunities for removing functionality
 - Opportunities for deferring init

Using initcall_debug

A terminal window titled "Terminal — vim — 66x19" showing the output of the command `journalctl | egrep 'calling|initcall' | egrep 'ubifs|ip_auto'`. The output shows two initcall sequences. The first sequence for `ubifs_init` shows a call time of 1 unit and a return time of 530 usecs. The second sequence for `ip_auto_config` shows a call time of 1 unit and a return time of 2777329 usecs. The values 530 and 2777329 are highlighted with green and red boxes respectively.

```
root@mx6q:~# journalctl | egrep 'calling|initcall' | egrep 'ubifs|ip_auto'
calling  ubifs_init+0x0/0xd0 @ 1
initcall ubifs_init+0x0/0xd0 returned 0 after 530 usecs
calling  ip_auto_config+0x0/0x1048 @ 1
initcall ip_auto_config+0x0/0x1048 returned 0 after 2777329 usecs
```

Oprofile

- Oprofile

- Statistical profiler which records PC of currently executing program when specified events occur
- Events include: (machine/processor dependent)
 - Timer, cache refills, interrupts while masked, many more
 - Varies by architecture
- Useful for finding “hotspots” in kernel and application code
 - Not available on every platform
 - Often requires some customization

Perf

- Much more functionality than oprofile
 - Generally lower overhead, no daemon required
- Statistical profiling of entire system
 - Kernel and userspace
- Integrates hardware performance counters, tracepoints and dynamic probes for advanced profiling
- Userspace app which controls operations
 - `perf`
- Several subcommands: `stat`, `top`, `record`, `report`, `sched...`
- Does not work on all platforms
 - May need porting/patching to your particular platform

Using strace to profile applications

- Strace can be used to collect timing information for a process
 - `strace -tt 2>/tmp/strace.log tthttpd ...`
- Determine where time is being spent in application startup
- Can also collect system call counts (-c)
- Can see time spent in each system call (-T)
- Great for finding extraneous operations (ubiquitous)
 - scanning invalid paths for files (e.g. dynamic libs, fonts, etc),
 - opening a file multiple times, etc.
- Strace can follow children (-o -ff)
- Strace adds **SIGNIFICANT** overhead to the execution of the program
 - Good for relative timings, not absolute
 - May slow execution so much that it “breaks” interaction with other processes



Optimization Techniques

Optimization Techniques

- There are numerous ways to speed up boot time
- Your mileage may vary depending on many factors
- Identify the longest bootup paths and select these for optimization
- Some techniques are obvious, others not so
- Some are aggressive and intrusive
- Many are simply tweaks and easy to apply
 - Turn off unneeded options, etc.

Optimizing U-Boot

- U-Boot must be relocated from Flash into DRAM
 - Reducing the image size reduces relocation time
- Lots of useful development functionality
 - tftp, pci scan, mem utils, disk utils, load*, dhcp, etc
 - In a production system, many of these features are unnecessary
 - Disabling these features can have a significant impact on boot time
- You want the bootloader to do it's work and get out of the way as fast as possible
- Look at every CONFIG_* option in your board configuration header
 - include/configs/<board_name>.h

Optimizing the Linux Kernel

- Size matters
 - The kernel needs to be loaded from FLASH into RAM
 - Smaller == faster
- Consider using an uncompressed kernel
 - Decompression can take dozens to hundreds of milliseconds
- Configure as many drivers as possible as modules
 - Mostly a "brute force" approach – trial and error, time consuming
- Consider using deferred initcall patch
 - Defer module init until much later in boot cycle
 - Initcalls deferred until triggered in userspace

Optimization Techniques

- XIP (Your mileage may vary)
- Limit console printk()
- Pre-configure or eliminate udev
- Kernel modules/deferred initcalls
- RTC_nosync
- Checkpoint restart
- Use parallelism for multi-core
- Cache systemd config
- Minimize rootfs size
- ...

Userspace – Optimize `init`

- Use BusyBox – Very popular in embedded systems
- Consider a custom `init` for very aggressive boot times
 - Can configure `init=myinit` on kernel command line
 - Allows complete customization of userland initialization
 - It is always faster to run native code than scripts
 - In general, every line of a script causes `fork()/exec()`
 - Often used in fixed function types of devices
- If you're using ready-made startup scripts
 - Eliminate unnecessary stuff (`set -x`)
 - Run multiple scripts in parallel wherever possible
 - May require adding some synchronization between services you start in parallel if there are dependencies.
- Use SystemD instead of SysV `init`

Using systemd

- Alternative to SysV Init
- Avoids much fork/exec of typical start up scripts
- Compatible with SysV Init scripts, but
 - Translate to systemd config files for best results
- Solves many of the startup dependencies quite nicely
- Parallelizes as much as it can
 - Big win on multicore
- See <http://0pointer.de/blog/projects/systemd.html> for an interesting introduction by the author

Application Prelink

- A good portion of application initialization time is spent resolving symbols to dynamic libraries
- Using Prelinking you can cut off a significant portion of application startup time if you have a large/complicated userland
- Tries to assign a preferred address space to each library used by an application – ahead of time
- Prelink is essential to rapid application startup

Designing your Applications - Considerations

- Keep it small
- Prelink
- Be careful adding dependencies on new libraries
 - it can snowball
- Keep fork/exec to minimum
- Some multithreading can help esp. with multicore
- Use Analysis tools – profiling, strace, etc.
- Use only Fast-path I/O Peripherals (e.g. – no Wifi)
- Avoid “discovery” code if possible



Achieving Sub-1 Second Boot

Aggressive Boot Time Reduction

- Mentor's Adaptive Preloading File System*
 - Profiles boot, stores read block list for faster playback
- U-Boot DMA
 - SPL or KL configuration
- Consider Read Only FS for root - YMMV
 - CRAMFS/SquashFS
- UBI Fastmap
- Avoid Initramfs
 - Apps that require shared libs require the entire lib in initramfs
 - Kernel drivers are more optimized than u-boot
 - Use a preloading FS instead
- Tiny kernel -> kexec full Linux kernel

*This is a technology, not a product

Linux Fast Boot – Mentor Projects

- Lesson Learned
 - Boot time is highly dependent on choices made in HW
 - Hardware architecture and system architecture matter
 - Individual requirements vary
- Mentor has significant experience
 - Kernel optimizations for custom boards in automotive space
 - Architecture analysis for system trade-offs in boot time performance
- Sample Benchmarks Mentor Graphics has met in the past
 - 100 msec boot loader to start kernel loading
 - 750 msec for Linux kernel start
 - *Rootfs and read/write filesystem mounted*
 - *First user process running*
 - Audio on within 1.25 sec
 - Safety-critical Camera feeds within 2 sec
 - Home screen available within 10 sec



**Thank you!
Your Questions
are Welcome**

mentor.com/embedded